

Statistics and Data Analysis

Unit 01 – Lecture 02: Feature Scaling and Feature Engineering Basics

Tofik Ali

School of Computer Science, UPES Dehradun

February 14, 2026

<https://github.com/tali7c/Statistics-and-Data-Analysis>

Overview
0

Feature Scaling
oooooooooooo

Transformations
oooo

Feature Engineering Basics
ooooooo

Demo
oo

Summary
o

Quick Links

Scaling

Transforms

Feature Engg.

Demo

Summary

Agenda

1 Overview

2 Feature Scaling

3 Transformations

4 Feature Engineering Basics

5 Demo

6 Summary

Learning Outcomes

- Explain why feature scaling is needed in many ML/analytics workflows

Learning Outcomes

- Explain why feature scaling is needed in many ML/analytics workflows
- Compute min–max normalization and z-score standardization

Learning Outcomes

- Explain why feature scaling is needed in many ML/analytics workflows
- Compute min–max normalization and z-score standardization
- Choose appropriate scaling/transformation for a given scenario

Learning Outcomes

- Explain why feature scaling is needed in many ML/analytics workflows
- Compute min–max normalization and z-score standardization
- Choose appropriate scaling/transformation for a given scenario
- Create basic engineered features (encoding, bins, interactions, date parts)

Why Scale Features?

- Different features can have very different units and ranges

Why Scale Features?

- Different features can have very different units and ranges
- Distance-based methods (kNN, k-means) are dominated by large-scale features

Why Scale Features?

- Different features can have very different units and ranges
- Distance-based methods (kNN, k-means) are dominated by large-scale features
- Optimization can be faster/more stable after scaling (gradient-based models)

Why Scale Features?

- Different features can have very different units and ranges
- Distance-based methods (kNN, k-means) are dominated by large-scale features
- Optimization can be faster/more stable after scaling (gradient-based models)
- Scaling also helps compare feature importance in some linear models

Min–Max Normalization

Maps a feature to [0, 1]:

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

- Keeps relative ordering

Min–Max Normalization

Maps a feature to [0, 1]:

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

- Keeps relative ordering
- Sensitive to extreme outliers (min/max can be unstable)

Standardization (z-score)

Centers and scales using mean and standard deviation:

$$z = \frac{x - \mu}{\sigma}$$

- Produces mean 0 and std 1 (approximately)

Standardization (z-score)

Centers and scales using mean and standard deviation:

$$z = \frac{x - \mu}{\sigma}$$

- Produces mean 0 and std 1 (approximately)
- Often a good default for many algorithms

Important: Avoid Data Leakage

- Scaling parameters (min/max, mean/std) must be computed on **training data only**

Important: Avoid Data Leakage

- Scaling parameters (min/max, mean/std) must be computed on **training data only**
- Then apply the same parameters to validation/test data

Important: Avoid Data Leakage

- Scaling parameters (min/max, mean/std) must be computed on **training data only**
- Then apply the same parameters to validation/test data
- Otherwise, test information leaks into training and results look unrealistically good

Exercise 1: Min–Max

Suppose a feature has $\min = 20$, $\max = 80$.

Compute the min–max normalized value for $x = 50$.

Solution 1

$$x' = \frac{50 - 20}{80 - 20} = \frac{30}{60} = 0.5$$

Exercise 2: z-score

A test score has mean $\mu = 60$ and standard deviation $\sigma = 10$.
Compute the z-score for $x = 80$ and interpret it.

Solution 2

$$z = \frac{80 - 60}{10} = 2$$

Interpretation: 80 is 2 standard deviations above the mean.

Exercise 3: Who Needs Scaling?

For each algorithm, answer: scaling important? (Yes/No)

- 1 kNN
- 2 k-means clustering
- 3 Linear regression (with gradient descent)
- 4 Decision tree

Solution 3

- kNN: Yes (distance-based)
- k-means: Yes (distance-based)
- Linear regression (GD): Often yes (helps optimization)
- Decision tree: Usually no (splits are scale-invariant)

Why Transform Features?

- Reduce skewness (e.g., income is often right-skewed)

Why Transform Features?

- Reduce skewness (e.g., income is often right-skewed)
- Stabilize variance and make patterns more linear

Why Transform Features?

- Reduce skewness (e.g., income is often right-skewed)
- Stabilize variance and make patterns more linear
- Improve interpretability and model fit in some cases

Log Transform (Common for Right-Skew)

For positive values, a common transform is:

$$x' = \log(1 + x)$$

- Compresses large values and spreads small values

Log Transform (Common for Right-Skew)

For positive values, a common transform is:

$$x' = \log(1 + x)$$

- Compresses large values and spreads small values
- Reduces the impact of extreme values

Exercise 4: Pick a Transform

Which is a reasonable first transform for a right-skewed income column?

- (a) x
- (b) x^2
- (c) $\log(1 + x)$

Solution 4

(c) $\log(1 + x)$ is a common choice for right-skewed positive features.

What is Feature Engineering?

Feature engineering means creating useful new variables from raw data to improve analysis/models.

- Encode categorical variables (one-hot)

What is Feature Engineering?

Feature engineering means creating useful new variables from raw data to improve analysis/models.

- Encode categorical variables (one-hot)
- Create buckets/bins (low/medium/high)

What is Feature Engineering?

Feature engineering means creating useful new variables from raw data to improve analysis/models.

- Encode categorical variables (one-hot)
- Create buckets/bins (low/medium/high)
- Interactions and ratios (effort = hours × attendance)

What is Feature Engineering?

Feature engineering means creating useful new variables from raw data to improve analysis/models.

- Encode categorical variables (one-hot)
- Create buckets/bins (low/medium/high)
- Interactions and ratios (effort = hours × attendance)
- Date parts (month, weekday)

What is Feature Engineering?

Feature engineering means creating useful new variables from raw data to improve analysis/models.

- Encode categorical variables (one-hot)
- Create buckets/bins (low/medium/high)
- Interactions and ratios ($\text{effort} = \text{hours} \times \text{attendance}$)
- Date parts (month, weekday)
- Simple text features (length, keywords)

Example (Student Dataset Columns)

Raw columns:

program, city, join_date, attendance_pct, study_hours_week,
cgpa, family_income_k, backlog

Possible engineered features:

- join_month, join_weekday

Example (Student Dataset Columns)

Raw columns:

program, city, join_date, attendance_pct, study_hours_week,
cgpa, family_income_k, backlog

Possible engineered features:

- join_month, join_weekday
- attendance_bucket (low/medium/high)

Example (Student Dataset Columns)

Raw columns:

```
program, city, join_date, attendance_pct, study_hours_week,  
           cgpa, family_income_k, backlog
```

Possible engineered features:

- join_month, join_weekday
- attendance_bucket (low/medium/high)
- income_log1p

Example (Student Dataset Columns)

Raw columns:

program, city, join_date, attendance_pct, study_hours_week,
cgpa, family_income_k, backlogs

Possible engineered features:

- join_month, join_weekday
- attendance_bucket (low/medium/high)
- income_log1p
- effort_index = hours × attendance/100

Example (Student Dataset Columns)

Raw columns:

```
program, city, join_date, attendance_pct, study_hours_week,  
        cgpa, family_income_k, backlogs
```

Possible engineered features:

- join_month, join_weekday
- attendance_bucket (low/medium/high)
- income_log1p
- effort_index = hours × attendance/100
- has_backlog (0/1)

Exercise 5: Design Features

You have: attendance_pct, study_hours_week, family_income_k, join_date, program.

Task: Propose **three** engineered features and explain why they might help.

Solution 5 (Examples)

- `income_log1p`: reduces right-skew and outlier impact
- `effort_index`: combines hours + attendance into one effort score
- `join_month`: captures seasonal patterns (admissions, attendance patterns)

Exercise 6: One-hot Encoding

Program has categories: CSE, ECE, AIML.

Task: Write the one-hot encoded columns for program.

Solution 6

One-hot columns:

- program_CSE
- program_ECE
- program_AIML

Each row has 1 for its category and 0 for the others.

Mini Demo (Python)

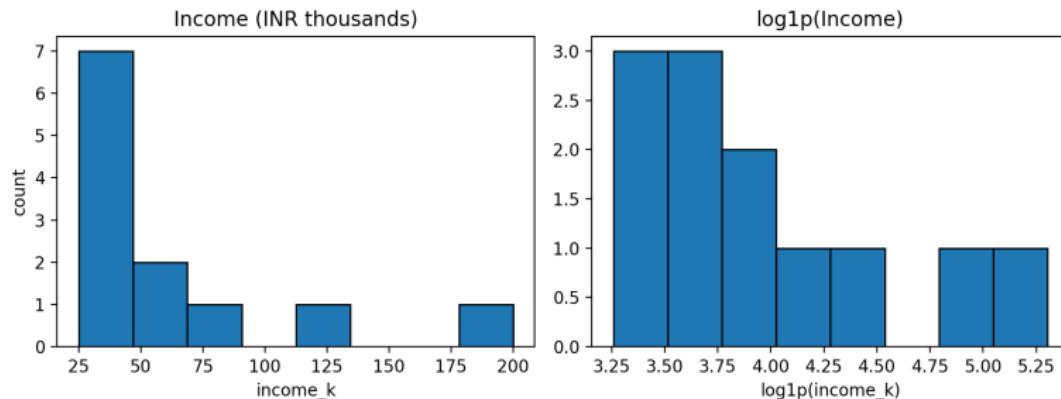
Run from the lecture folder:

```
python demo/scaling_feature_engineering_demo.py
```

Outputs:

- `data/student_features_engineered.csv`
- plots in `images/` (income vs log-income, hours vs CGPA)

Demo Output (Example Plot)



Summary

- Scaling makes features comparable and prevents domination by large-scale variables

Exit question: Why is scaling “fit on train only” and then applied to test?

Summary

- Scaling makes features comparable and prevents domination by large-scale variables
- Min–max maps to $[0, 1]$; z-score centers and scales by mean/std

Exit question: Why is scaling “fit on train only” and then applied to test?

Summary

- Scaling makes features comparable and prevents domination by large-scale variables
- Min–max maps to $[0, 1]$; z-score centers and scales by mean/std
- Transformations like $\log(1 + x)$ help reduce skew and outlier influence

Exit question: Why is scaling “fit on train only” and then applied to test?

Summary

- Scaling makes features comparable and prevents domination by large-scale variables
- Min–max maps to $[0, 1]$; z-score centers and scales by mean/std
- Transformations like $\log(1 + x)$ help reduce skew and outlier influence
- Feature engineering creates informative variables (encoding, bins, interactions, dates)

Exit question: Why is scaling “fit on train only” and then applied to test?